

Big Mart Sales

1. Loading important files

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
        5 %matplotlib inline
```

2. Loading Dataset

```
In [2]: 1 df_train=pd.read_csv(r'D:\New folder\New folder (2)\bigmart\Train.csv')
        2 df_test =pd.read_csv(r'D:\New folder\New folder (2)\bigmart\Test.csv')
```

```
In [3]: 1 df_train.head()
```

```
Out[3]:
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Type
0	FDA15	9.30	Low Fat	0.016047	Dairy	249.8092	OUT049	1999	M
1	DRC01	5.92	Regular	0.019278	Soft Drinks	48.2692	OUT018	2009	M
2	FDN15	17.50	Low Fat	0.016760	Meat	141.6180	OUT049	1999	M
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	OUT010	1998	M
4	NCD19	8.93	Low Fat	0.000000	Household	53.8614	OUT013	1987	M

In [4]: 1 df_test.head()

Out[4]:

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet
0	FDW58	20.750	Low Fat	0.007565	Snack Foods	107.8622	OUT049	1999	M
1	FDW14	8.300	reg	0.038428	Dairy	87.3198	OUT017	2007	
2	NCN55	14.600	Low Fat	0.099575	Others	241.7538	OUT010	1998	
3	FDQ58	7.315	Low Fat	0.015388	Snack Foods	155.0340	OUT017	2007	
4	FDY38	NaN	Regular	0.118599	Dairy	234.2300	OUT027	1985	M

. Checking Information

In [5]: 1 df_train.shape

Out[5]: (8523, 12)

In [6]: 1 df_test.shape

Out[6]: (5681, 11)

Checking Describe

In [7]: 1 df_train.describe()

Out[7]:

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	7060.000000	8523.000000	8523.000000	8523.000000	8523.000000
mean	12.857645	0.066132	140.992782	1997.831867	2181.288914
std	4.643456	0.051598	62.275067	8.371760	1706.499616
min	4.555000	0.000000	31.290000	1985.000000	33.290000
25%	8.773750	0.026989	93.826500	1987.000000	834.247400
50%	12.600000	0.053931	143.012800	1999.000000	1794.331000
75%	16.850000	0.094585	185.643700	2004.000000	3101.296400
max	21.350000	0.328391	266.888400	2009.000000	13086.964800

In [8]: 1 df_test.describe()

Out[8]:

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year
count	4705.000000	5681.000000	5681.000000	5681.000000
mean	12.695633	0.065684	141.023273	1997.828903
std	4.664849	0.051252	61.809091	8.372256
min	4.555000	0.000000	31.990000	1985.000000
25%	8.645000	0.027047	94.412000	1987.000000
50%	12.500000	0.054154	141.415400	1999.000000
75%	16.700000	0.093463	186.026600	2004.000000
max	21.350000	0.323637	266.588400	2009.000000

Checking Info

In [9]: 1 df_train.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8523 entries, 0 to 8522
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Item_Identifier                       8523 non-null   object
1   Item_Weight                          7060 non-null   float64
2   Item_Fat_Content                     8523 non-null   object
3   Item_Visibility                      8523 non-null   float64
4   Item_Type                           8523 non-null   object
5   Item_MRP                            8523 non-null   float64
6   Outlet_Identifier                    8523 non-null   object
7   Outlet_Establishment_Year            8523 non-null   int64
8   Outlet_Size                          6113 non-null   object
9   Outlet_Location_Type                 8523 non-null   object
10  Outlet_Type                          8523 non-null   object
11  Item_Outlet_Sales                    8523 non-null   float64
dtypes: float64(4), int64(1), object(7)
memory usage: 799.2+ KB
```

In [10]: 1 df_test.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5681 entries, 0 to 5680
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Item_Identifier        5681 non-null  object
1   Item_Weight            4705 non-null  float64
2   Item_Fat_Content       5681 non-null  object
3   Item_Visibility        5681 non-null  float64
4   Item_Type              5681 non-null  object
5   Item_MRP               5681 non-null  float64
6   Outlet_Identifier      5681 non-null  object
7   Outlet_Establishment_Year 5681 non-null  int64
8   Outlet_Size            4075 non-null  object
9   Outlet_Location_Type   5681 non-null  object
10  Outlet_Type            5681 non-null  object
dtypes: float64(3), int64(1), object(7)
memory usage: 488.3+ KB
```

. Checking missing values

In [11]: 1 df_train.isnull().sum()

```
Out[11]: Item_Identifier        0
Item_Weight            1463
Item_Fat_Content       0
Item_Visibility        0
Item_Type              0
Item_MRP               0
Outlet_Identifier      0
Outlet_Establishment_Year 0
Outlet_Size            2410
Outlet_Location_Type   0
Outlet_Type            0
Item_Outlet_Sales      0
dtype: int64
```

```
In [12]: 1 df_test.isnull().sum()
```

```
Out[12]: Item_Identifier      0
Item_Weight      976
Item_Fat_Content      0
Item_Visibility      0
Item_Type      0
Item_MRP      0
Outlet_Identifier      0
Outlet_Establishment_Year      0
Outlet_Size      1606
Outlet_Location_Type      0
Outlet_Type      0
dtype: int64
```

Filling Missing Values

```
In [13]: 1 # 'Item_Weight' has a numerical values so will fill it by using mean imputation
2 df_train['Item_Weight'].fillna(df_train['Item_Weight'].mean(),inplace=True)
3 df_test['Item_Weight'].fillna(df_test['Item_Weight'].mean(),inplace=True)
```

```
In [14]: 1 df_train['Item_Weight'].mode()
```

```
Out[14]: 0    12.857645
dtype: float64
```

```
In [15]: 1 # Outlet_Size has a categorical values so we will fill it by using mode imputation
2 df_test['Outlet_Size'].mode()
```

```
Out[15]: 0    Medium
dtype: object
```

```
In [16]: 1 df_train['Outlet_Size'].fillna(df_train['Outlet_Size'].mode()[0],inplace=True)
```

```
In [17]: 1 df_test['Outlet_Size'].fillna(df_test['Outlet_Size'].mode()[0],inplace=True)
```

```
In [18]: 1 df_train.isnull().sum()
```

```
Out[18]: Item_Identifier      0
Item_Weight      0
Item_Fat_Content  0
Item_Visibility  0
Item_Type      0
Item_MRP      0
Outlet_Identifier  0
Outlet_Establishment_Year  0
Outlet_Size      0
Outlet_Location_Type  0
Outlet_Type      0
Item_Outlet_Sales  0
dtype: int64
```

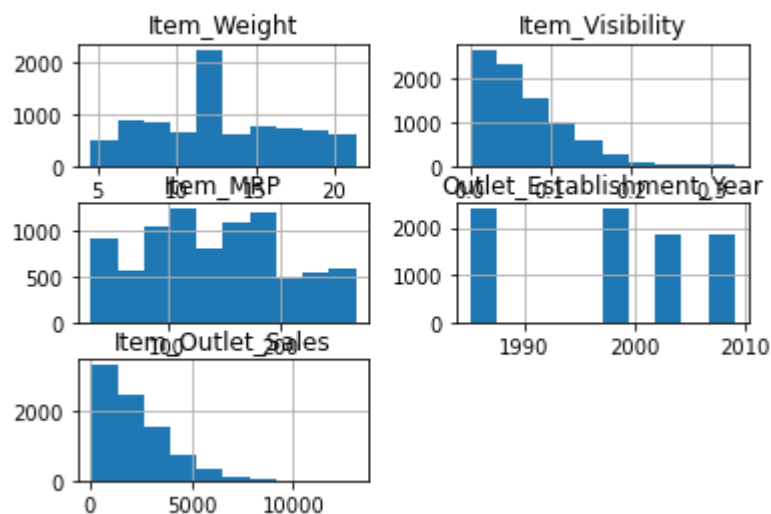
Data Reduction

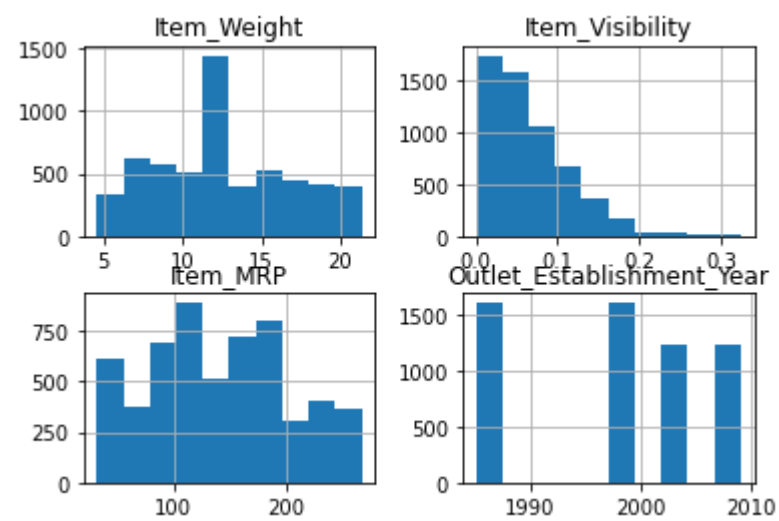
```
In [19]: 1 # We dont need the two columns i.e, ['Item_Identifier'] ,['Outlet_Identifier'] so we will drop this columns
2 df_train.drop(columns=['Item_Identifier', 'Outlet_Identifier'], inplace=True)
3 df_test.drop(columns=['Item_Identifier', 'Outlet_Identifier'], inplace=True)
```

EDA(Exploratory Data Analysis)


```
In [20]: 1 df_train.hist() ,df_test.hist()
```

```
Out[20]: (array([[<AxesSubplot:title={'center':'Item_Weight'}>,
<AxesSubplot:title={'center':'Item_Visibility'}>],
[<AxesSubplot:title={'center':'Item_MRP'}>,
<AxesSubplot:title={'center':'Outlet_Establishment_Year'}>],
[<AxesSubplot:title={'center':'Item_Outlet_Sales'}>,
<AxesSubplot:>]], dtype=object),
array([[<AxesSubplot:title={'center':'Item_Weight'}>,
<AxesSubplot:title={'center':'Item_Visibility'}>],
[<AxesSubplot:title={'center':'Item_MRP'}>,
<AxesSubplot:title={'center':'Outlet_Establishment_Year'}>]],
dtype=object))
```





```
In [21]: 1 plt.figure(figsize =(10,5))
2         sns.heatmap(df_train.corr(),annot=True)
3         plt.show()
```



```
In [22]: 1 plt.figure(figsize=(10,5))
2         sns.heatmap(df_test.corr(),annot=True)
3         plt.show()
```



```
In [23]: 1 df_train.head()
```

```
Out[23]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type
0	9.30	Low Fat	0.016047	Dairy	249.8092	1999	Medium	Tier 1	Supermarket
1	5.92	Regular	0.019278	Soft Drinks	48.2692	2009	Medium	Tier 3	Supermarket
2	17.50	Low Fat	0.016760	Meat	141.6180	1999	Medium	Tier 1	Supermarket
3	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	1998	Medium	Tier 3	Supermarket
4	8.93	Low Fat	0.000000	Household	53.8614	1987	High	Tier 3	Supermarket

Preprocessing task before model building

Label Encoding

```
In [24]: 1 from sklearn.preprocessing import LabelEncoder
        2 le=LabelEncoder()
```

```
In [25]: 1 df_train['Item_Fat_Content'] = le.fit_transform(df_train['Item_Fat_Content'])
        2 df_train['Item_Type'] = le.fit_transform(df_train['Item_Type'])
        3 df_train['Outlet_Size'] = le.fit_transform(df_train['Outlet_Size'])
        4 df_train['Outlet_Type'] = le.fit_transform(df_train['Outlet_Type'])
        5 df_train['Outlet_Location_Type'] = le.fit_transform(df_train['Outlet_Location_Type'])
```

```
In [26]: 1 df_train
```

```
Out[26]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type
0	9.300	1	0.016047	4	249.8092	1999	1	0
1	5.920	2	0.019278	14	48.2692	2009	1	2
2	17.500	1	0.016760	10	141.6180	1999	1	0
3	19.200	2	0.000000	6	182.0950	1998	1	2
4	8.930	1	0.000000	9	53.8614	1987	0	2
...
8518	6.865	1	0.056783	13	214.5218	1987	0	2
8519	8.380	2	0.046982	0	108.1570	2002	1	1
8520	10.600	1	0.035186	8	85.1224	2004	2	1
8521	7.210	2	0.145221	13	103.1332	2009	1	2
8522	14.800	1	0.044878	14	75.4670	1997	2	0

8523 rows × 10 columns



Splitting our data into train and test

```
In [27]: 1 x=df_train.drop(['Item_Outlet_Sales'],axis=1)
        2 y=df_train['Item_Outlet_Sales']
```

```
In [28]: 1 from sklearn.model_selection import train_test_split
```

```
In [29]: 1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=101)
```

Standarisation

```
In [30]: 1 x.describe()
```

```
Out[30]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type
count	8523.000000	8523.000000	8523.000000	8523.000000	8523.000000	8523.000000	8523.000000	8523.000000
mean	12.857645	1.369354	0.066132	7.226681	140.992782	1997.831867	1.170832	1.111111
std	4.226124	0.644810	0.051598	4.209990	62.275067	8.371760	0.600327	0.811111
min	4.555000	0.000000	0.000000	0.000000	31.290000	1985.000000	0.000000	0.000000
25%	9.310000	1.000000	0.026989	4.000000	93.826500	1987.000000	1.000000	0.000000
50%	12.857645	1.000000	0.053931	6.000000	143.012800	1999.000000	1.000000	1.000000
75%	16.000000	2.000000	0.094585	10.000000	185.643700	2004.000000	2.000000	2.000000
max	21.350000	4.000000	0.328391	15.000000	266.888400	2009.000000	2.000000	2.000000

```
In [31]: 1 # Here our data are not close like Item_Weight is very high where Outlet_Location_Type is very small. To over
```

```
In [32]: 1 from sklearn.preprocessing import StandardScaler
        2 sc=StandardScaler()
```

```
In [33]: 1 x_train_std=sc.fit_transform(x_train)
```

```
In [34]: 1 x_test_std=sc.fit_transform(x_test)
```

```
In [35]: 1 x_train_std
```

```
Out[35]: array([[ 1.52290029, -0.57382672,  0.68469729, ..., -1.95699503,
                  1.08786619, -0.25964107],
                [-1.23985603, -0.57382672, -0.09514748, ..., -0.28872895,
                  -0.13870429, -0.25964107],
                [ 1.54667616,  0.97378032, -0.00838589, ..., -0.28872895,
                  -0.13870429, -0.25964107],
                ...,
                [-0.08197107, -0.57382672, -0.9191623 , ...,  1.37953713,
                  -1.36527477, -0.25964107],
                [-0.74888428,  0.97378032,  1.21363058, ..., -0.28872895,
                  -0.13870429, -0.25964107],
                [ 0.67885683, -0.57382672,  1.83915356, ..., -0.28872895,
                  1.08786619,  0.98524841]])
```

```
In [36]: 1 y_train
```

```
Out[36]: 3684      163.7868
          1935      1607.2412
          5142      1510.0344
          4978      1784.3440
          2299      3558.0352
          ...
          599      5502.8370
          5695      1436.7964
          8006      2167.8448
          1361      2700.4848
          1547      829.5868
          Name: Item_Outlet_Sales, Length: 6818, dtype: float64
```

```
In [37]: 1 y_test
```

```
Out[37]: 8179      904.8222
          8355      2795.6942
          3411      1947.4650
          7089      872.8638
          6954      2450.1440
          ...
          1317      1721.0930
          4996      914.8092
          531       370.1848
          3891      1358.2320
          6629      2418.1856
          Name: Item_Outlet_Sales, Length: 1705, dtype: float64
```

```
In [38]: 1 import joblib
```

```
In [ ]: 1
```

Model Building

```
In [39]: 1 # Using Alogithm of Linear Regresion
```

```
In [40]: 1 from sklearn.linear_model import LinearRegression
          2 lr= LinearRegression()
```

```
In [41]: 1 lr.fit(x_train_std,y_train)
```

```
Out[41]: LinearRegression()
```

```
In [42]: 1 y_pred_lr=lr.predict(x_test_std)
```

```
In [43]: 1 from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
```

```
In [44]: 1 print(r2_score(y_test,y_pred_lr))
          2 print( mean_absolute_error(y_test,y_pred_lr))
          3 print(np.sqrt(mean_squared_error(y_test,y_pred_lr)))
```

0.5020054018117118

885.7810720872159

1164.9965315232587

```
In [45]: 1 # Using Alogithm of Random Forest Regresion
```

```
In [ ]: 1 from sklearn.ensemble import RandomForestRegressor
          2 rf=RandomForestRegressor(n_estimators=1000) #wrote after finding the valuye of n_estimator by applying
          3                                             #hyperparameter tuning
```

```
In [ ]: 1 rf.fit(x_train,y_train)
```

```
In [ ]: 1 y_pred_rf=rf.predict(x_test)
```

```
In [ ]: 1 print(r2_score(y_test,y_pred_rf))
          2 print( mean_absolute_error(y_test,y_pred_rf))
          3 print(np.sqrt(mean_squared_error(y_test,y_pred_rf)))
```

Hyperparameter tuning


```
In [50]: 1 from sklearn.model_selection import GridSearchCV, RepeatedStratifiedKFold
2 #define model and parameter
3 model=RandomForestRegressor()
4 n_estimators=[10,100,1000]
5 max_depth=range(1,31)
6 min_sample_leaf=np.linspace(0.1,1.0)
7 max_features=['auto','sqrt','log2']
8
9 #define grid search
10 grid=dict(n_estimators=n_estimators)
11 #cv=RepeatedStratifiedKFold(n_splits=5,n_repeats=3, random_state=101)
12 grid_search_forest=GridSearchCV(estimator=model,param_grid=grid,n_jobs=-1,
13                                 scoring='r2', error_score=0, verbose=2, cv=2)
14 grid_search_forest.fit(x_train_std,y_train)
15 # Summerize result
16 print(f'Best:{grid_search_forest.best_score_: .3f}using{grid_search_forest.best_params_}')
17 means=grid_search_forest.cv_results_['mean_test_score']
18 stds=grid_search_forest.cv_results_['std_test_score']
19 params=grid_search_forest.cv_results_['params']
20
21 for mean,stdev, param in zip(means,stds,params):
22     print(f'{mean: .3f}({stdev: .3f})with{param}')
```

Fitting 2 folds for each of 3 candidates, totalling 6 fits

```
Best:0.550using{'n_estimators': 1000}
0.517(0.011)with{'n_estimators': 10}
0.542(0.007)with{'n_estimators': 100}
0.550(0.005)with{'n_estimators': 1000}
```

```
In [51]: 1 grid_search_forest.best_params_
```

```
Out[51]: {'n_estimators': 1000}
```

```
In [52]: 1 grid_search_forest.best_score_
```

```
Out[52]: 0.5498216664771529
```

```
In [53]: 1 y_pred_rf_grid= grid_search_forest.predict(x_test_std)
```

In [54]: 1 r2_score(y_test,y_pred_rf_grid)

Out[54]: 0.5506195416916224

In []:

1

By: Ashwin Gorakhnath Dubey

In []:

1